

An Introduction to the Tidyverse

David Zimmermann
R User Group Oxford
02 October 2017



Outline

1. An Overview of the Tidyverse
2. Pipes %>%
3. Tidy Data and tidyr
4. Data Munging and dplyr

An Overview of the Tidyverse

A collection of libraries for data science with a consistent design and API

- **tibble** data.frame extension
- **tidyr** reshaping data
- **dplyr** data munging and manipulation
- readr read and write to text-files (csv, txt, etc)
- ggplot2 data visualization
- purrr functional programming
- ...

Lives here <https://www.tidyverse.org/> and here <https://github.com/tidyverse/tidyverse>

Usage

Install once

```
install.packages("tidyverse")
```

Usage in Scripts

```
library(tidyverse)
```

```
# and you are good to go
```

```
df <- data_frame(id = 1:3, name = c("Alice", "Bob", "Charlie"))
```

```
df
```

```
# A tibble: 3 x 2
```

```
  id  name  
  <int> <chr>
```

```
1    1 Alice
```

```
2    2  Bob
```

```
3    3 Charlie
```

data?frame

```
df_dot      <- data.frame(id = 1:2, name = c("Alice", "Bob"))  
df_underscore <- data_frame(id = 1:2, name = c("Alice", "Bob"))
```

```
df_dot
```

```
  id name  
1  1 Alice  
2  2  Bob
```

```
df_underscore
```

```
# A tibble: 2 x 2  
  id name  
  <int> <chr>  
1     1 Alice  
2     2  Bob
```

data?frame cont'd

`data_frame()/tibble()`:

- doesn't convert strings to factors and allows for lists as elements!
- prints more information in a nicer format (try with a dataset of >10 columns/rows)
- is still a `data.frame`, thus works with older code!

```
str(df_dot)
```

```
'data.frame': 2 obs. of 2 variables:  
 $ id : int 1 2  
 $ name: Factor w/ 2 levels "Alice","Bob": 1 2
```

```
str(df_underscore)
```

```
Classes 'tbl_df', 'tbl' and 'data.frame': 2 obs. of 2 variables:
```



Ceci n'est pas une pipe.

Magnette

Pipes

Then:

```
value2 <- foo(value)  
value3 <- bar(value2)
```

```
# Or  
bar(foo(value))
```


Pipes

Then:

```
value2 <- foo(value)
value3 <- bar(value2)

# Or
bar(foo(value))
```

Take the output (value) of one function and use it in another function.

Now using the pipe-operator `%>%` (read it as “then”):

```
value %>% foo() %>% bar()
# or with better formatting
value %>%
  foo() %>%
  bar()
```

Pipes cont'd

In general

```
f(g(h(x)))
```

```
# becomes
```

```
x %>% h() %>% g() %>% f()
```

Pipes cont'd

In general

```
f(g(h(x)))  
# becomes  
x %>% h() %>% g() %>% f()
```

Example

Non-pipe

```
flights_dec <- filter(flights, month == 12)  
flights_dec_grouped <- group_by(flights_dec, day)  
summarise(flights_dec_grouped, mean_delay = mean(arr_delay))
```

vs. pipe

```
flights %>%  
  filter(month == 12) %>%  
  group_by(day) %>%  
  summarise(mean_delay = mean(arr_delay))
```

Pipes cont'd

Named arguments and arg-numbers using the `.`-argument

```
foo(w, x)
# becomes
x %>% foo(w, .)

foo(w, x, y, z)
# becomes
x %>% foo(w, ., y, z)

# Named Arguments
x %>% foo(w, x = .)
```

More information: <http://r4ds.had.co.nz/pipes.html>



Tidy Data

How should we store data in a “good” format?

Say, we have sales data for

- Alice
- Bob
- Charlie

For the years 2010 and 2011

Tidy Data cont'd

Wide-format

Year	Alice	Bob	Charlie
2010	105	100	90
2011	110	97	95

Long-format

Name	Year	Sales
Alice	2010	105
Alice	2011	110
Bob	2010	100
Bob	2011	97
Charlie	2010	90
Charlie	2011	95

Tidy Data cont'd

Wide-format

Year	Alice	Bob	Charlie
2010	105	100	90
2011	110	97	95

Long/Tidy-Data

Name	Year	Sales
Alice	2010	105
Alice	2011	110
Bob	2010	100
Bob	2011	97
Charlie	2010	90
Charlie	2011	95

Ease of use vs. size of data?

What happens if we receive data for 2012?

What happens if Dave joins?

What happens if we receive the number of visits per salesperson?

What happens if Alice and Charlie belong to Company X, but Bob to Company Y?

Formalised Tidy Data

One Variable per **Column**

One Observation per **Row**

country	year	cases	population
Afghanistan	1999	745	19987071
Afghanistan	2000	2666	20593360
Brazil	1999	37737	172006362
Brazil	2000	80488	174604898
China	1999	212258	1272915272
China	2000	216766	1280425583

variables

country	year	cases	population
Afghanistan	1999	745	19987071
Afghanistan	2000	2666	20593360
Brazil	1999	37737	172006362
Brazil	2000	80488	174604898
China	1999	212258	1272915272
China	2000	216766	1280425583


observations

country	year	cases	population
Afghanistan	1999	745	19987071
Afghanistan	2000	2666	20593360
Brazil	1999	37737	172006362
Brazil	2000	80488	174604898
China	1999	212258	1272915272
China	2000	216766	1280425583

values

Reshaping Data using TidyR (Tidyverse)

Wide-Data




`tidyr::gather()`



`tidyr::spread()`

Long/Tidy-Data



Reshaping Data

Recreate the Data

```
wide <- data_frame(year = c(2010, 2011),  
                    Alice = c(105, 110),  
                    Bob = c(100, 97),  
                    Charlie = c(90, 95))
```

wide

```
# A tibble: 2 x 4  
  year Alice Bob Charlie  
  <dbl> <dbl> <dbl> <dbl>  
1  2010  105  100    90  
2  2011  110   97    95
```

Reshaping Data Wide to Long

gather() the data to a long/tidy-format

```
long <- wide %>% gather(key = "name", value = "sales", -year)
long
```

```
# A tibble: 6 x 3
  year  name sales
  <dbl> <chr> <dbl>
1 2010  Alice  105
2 2011  Alice  110
3 2010   Bob   100
4 2011   Bob   97
5 2010 Charlie  90
6 2011 Charlie  95
```

- **key**: the name of the (future) variable that holds the key (the values in the header)
- **value**: the name of the (future) variable that holds the values (in the body)

Reshaping Data Long to Wide

spread() the data to a wide-format


```
wide2 <- long %>% spread(key = "name", value = "sales")
wide2
```

```
# A tibble: 2 x 4
  year Alice Bob Charlie
*   <dbl> <dbl> <dbl>   <dbl>
1  2010  105  100     90
2  2011  110   97     95
```

- **key**: the name of the variable that will go in the header
- **value**: the name of the variable that will go to the body

```
identical(wide, wide2)
```

```
[1] TRUE
```



www.rstudio.com

dplyr Overview

Grammar of Data Manipulation

Reasonably fast and consistent API

- `filter()` filter observations / rows
- `arrange()` arrange (sort) the dataset by a column
- `select()` select variables / columns
- `mutate()` change or create a variable
- `summarise()` summarise the dataset to a single row
- + `group_by()` operate by a grouping-variable

Lives here <http://dplyr.tidyverse.org/> and here
<https://github.com/tidyverse/dplyr>

Dataset

```
# install.packages("nycflights13")  
library(nycflights13)  
flights %>% glimpse()
```

Observations: 336,776

Variables: 19

```
$ year      <int> 2013, 2013, 2013, 2013, 2013, 2013, ...  
$ month     <int> 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, ...  
$ day       <int> 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, ...  
$ dep_time  <int> 517, 533, 542, 544, 554, 554, 555, ...  
$ sched_dep_time <int> 515, 529, 540, 545, 600, 558, 600, ...  
$ dep_delay <dbl> 2, 4, 2, -1, -6, -4, -5, -3, -3, -2...  
$ arr_time  <int> 830, 850, 923, 1004, 812, 740, 913, ...  
$ sched_arr_time <int> 819, 830, 850, 1022, 837, 728, 854, ...  
$ arr_delay <dbl> 11, 20, 33, -18, -25, 12, 19, -14, ...  
$ carrier   <chr> "UA", "UA", "AA", "B6", "DL", "UA", ...  
$ flight    <int> 1545, 1714, 1141, 725, 461, 1696, 5...  
$ tailnum   <chr> "N14228", "N24211", "N619AA", "N804...  
$ origin    <chr> "EWR", "LGA", "JFK", "JFK", "LGA", ...  
$ dest      <chr> "IAH", "IAH", "MIA", "BQN", "ATL", ...  
$ air_time  <dbl> 227, 227, 160, 183, 116, 150, 158, ...  
$ distance  <dbl> 1400, 1416, 1089, 1576, 762, 719, 1...  
$ hour      <dbl> 5, 5, 5, 5, 6, 5, 6, 6, 6, 6, 6, ...  
$ minute    <dbl> 15, 29, 40, 45, 0, 58, 0, 0, 0, 0, ...
```


Filter Observations

filter() the rows of a dataset for certain values only

Input Data

var		
x		
y		
z		
x		
z		
x		

```
df %>%  
  filter(var == "x")
```

Output Data

var		
x		
x		
x		

Filter Observations cont'd

Query: Find the long-distance flights in the spring of 2013.

```
flights %>%  
  filter(month <= 03 & distance > 2500)
```

```
# A tibble: 2,916 x 19
```

	year	month	day	dep_time	sched_dep_time	dep_delay
	<int>	<int>	<int>	<int>	<int>	<dbl>
1	2013	1	1	558	600	-2
2	2013	1	1	611	600	11
3	2013	1	1	655	700	-5
4	2013	1	1	729	730	-1
5	2013	1	1	734	737	-3
6	2013	1	1	745	745	0
7	2013	1	1	746	746	0
8	2013	1	1	803	800	3
9	2013	1	1	826	817	9
10	2013	1	1	857	900	-3

```
# ... with 2,906 more rows, and 13 more variables:
```

```
#   arr_time <int>, sched_arr_time <int>, arr_delay <dbl>,  
#   carrier <chr>, flight <int>, tailnum <chr>, origin <chr>,  
#   dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,  
#   minute <dbl>, time_hour <dtm>
```

Arrange Observations

arrange() the rows of a dataset for certain variables / columns

Input Data

[illegible]

```
df %>%  
  arrange(var1,  
           -var2)
```

Output Data

[illegible]

Arrange Observations cont'd

Query: Sort the flights by day (ascending) and departure-delay (descending).

```
flights %>%  
  arrange(day, -dep_delay)
```

```
# A tibble: 336,776 x 19  
  year month   day dep_time sched_dep_time dep_delay  
  <int> <int> <int>   <int>         <int>         <dbl>  
1  2013     1     1    848             1835          853  
2  2013    12     1    657             1930          687  
3  2013     5     1     9             1655          434  
4  2013     1     1   2343             1724          379  
5  2013     8     1   2311             1659          372  
6  2013     3     1   1528              920          368  
7  2013     7     1   1602              959          363  
8  2013     3     1   1449              855          354  
9  2013     7     1   2118             1525          353  
10 2013     7     1   1410              820          350  
# ... with 336,766 more rows, and 13 more variables:  
#   arr_time <int>, sched_arr_time <int>, arr_delay <dbl>,  
#   carrier <chr>, flight <int>, tailnum <chr>, origin <chr>,  
#   dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>
```

Select Variables

select() certain columns of a dataset

Input Data

var1	var2	var3

```
df %>%  
  select(var3,  
         foo = var1)
```

Output Data

var3	foo

Select Variables cont'd

Query: Select the carrier and the tail-number of the flights.

```
flights %>%  
  select(carrier, tail_number = tailnum)
```

```
# A tibble: 336,776 x 2
```

```
  carrier tail_number
```

```
    <chr>    <chr>
```

```
1     UA    N14228
```

```
2     UA    N24211
```

```
3     AA    N619AA
```

```
4     B6    N804JB
```

```
5     DL    N668DN
```

```
6     UA    N39463
```

```
7     B6    N516JB
```

```
8     EV    N829AS
```

```
9     B6    N593JB
```

```
10    AA    N3ALAA
```

```
# ... with 336,766 more rows
```

Mutate Variables

mutate() (change or create new) variables of a dataset

Input Data

var	Blue	Blue
Red	Green	Green
Light Orange	Green	Green
White	Green	Green
Light Orange	Green	Green
Red	Green	Green
White	Green	Green

```
df %>%
  mutate(var = var * 2,
         foo = 1:n())
```

Output Data

[illegible]

Mutate Variables cont'd

Query: Create a unique ID for each flight and compute the log distance.

```
flights %>%  
  mutate(id = 1:n(), log_dist = log(distance)) %>%  
  select(id, log_dist)
```

```
# A tibble: 336,776 x 2
```

	id	log_dist
	<int>	<dbl>
1	1	7.244228
2	2	7.255591
3	3	6.993015
4	4	7.362645
5	5	6.635947
6	6	6.577861
7	7	6.970730
8	8	5.433722
9	9	6.850126
10	10	6.597146

```
# ... with 336,766 more rows
```


Summarise Variables

summarise() (compute a summary of) variables of a dataset

Input Data

var1	var2	

```
df %>%  
  summarise(  
    mu_v1 = mean(var1),  
    min_v2 = min(var2),  
    max_v2 = max(var2),  
  )
```

Output Data

mu_v1	min_v2	max_v2

Summarise Variables cont'd

Query: Find the minimum, average, and maximum arrival delay for all flights.

```
flights %>%  
  filter(!is.na(arr_delay)) %>%  
  summarise(min_delay = min(arr_delay),  
            avg_delay = mean(arr_delay),  
            max_delay = max(arr_delay))
```

```
# A tibble: 1 x 3  
  min_delay avg_delay max_delay  
    <dbl>    <dbl>    <dbl>  
1    -86  6.895377    1272
```

Group Mutate

group_by() mutate (change or create new) variables of a dataset per group

Input Data

v	grp	

```
df %>%  
  group_by(grp) %>%  
  mutate(v = mean(v),  
         foo = 1:n())
```

Output Data

v	grp		foo

The data output is still grouped, to ungroup use

```
df %>% ... %>% ungroup()
```

Group Mutate cont'd

Query: For each flight, find the difference of the arrival-delay to the average arrival-delay of the respective carrier (airline).

```
flights %>%  
  filter(!is.na(arr_delay)) %>%  
  group_by(carrier) %>%  
  mutate(delta_arr_delay = arr_delay - mean(arr_delay)) %>%  
  select(delta_arr_delay)
```

```
# A tibble: 327,346 x 2  
# Groups:   carrier [16]  
  carrier delta_arr_delay  
    <chr>         <dbl>  
1    UA          7.441989  
2    UA         16.441989  
3    AA         32.635709  
4    B6        -27.457973  
5    DL        -26.644341  
6    UA          8.441989  
7    B6          9.542027  
8    EV        -29.796431  
9    B6        -17.457973  
10   AA          7.635709
```

Group Summarise

group_by() summarise (compute a summary of) variables of a dataset per group

Input Data

var1	var2	grp

```
df %>%  
  group_by(grp) %>%  
  summarise(  
    mu_v1 = mean(var1),  
    min_v2 = min(var2),  
    max_v2 = max(var2),  
  )
```

Output Data

grp	mu_v1	min_v2	max_v2

Group Summarise cont'd

Query: For each carrier (airline), find the mean and the median arrival delay over all flights in 2013.

```
flights %>%  
  filter(!is.na(arr_delay)) %>%  
  group_by(carrier) %>%  
  summarise(mean_delay = mean(arr_delay),  
            median_delay = median(arr_delay))
```

```
# A tibble: 16 x 3  
  carrier mean_delay median_delay  
  <chr>    <dbl>      <dbl>  
1     9E  7.3796692      -7  
2     AA  0.3642909      -9  
3     AS -9.9308886     -17  
4     B6  9.4579733       -3  
5     DL  1.6443409       -8  
6     EV 15.7964311       -1  
7     F9 21.9207048        6  
8     FL 20.1159055        5  
9     HA -6.9152047     -13  
10    MQ 10.7747334       -1  
#   with 6 more rows
```

Full Pipeline Example

Query: Find the 5 aircrafts (by tail number) that have the most time made-up (on average) and have at least 20 flights.

Full Pipeline Example

Query: Find the 5 aircrafts (by tail number) that have the most time made-up (on average) and have at least 20 flights.

```
flights %>%  
  filter(!is.na(arr_delay) & !is.na(dep_delay)) %>%  
  mutate(time_made_up = dep_delay - arr_delay) %>%  
  group_by(tailnum) %>%  
  summarise(n_flights = n(),  
            time_made_up = mean(time_made_up)) %>%  
  filter(n_flights > 20) %>%  
  arrange(-time_made_up) %>%  
  top_n(5)
```



```
# A tibble: 5 x 3  
  tailnum n_flights time_made_up  
  <chr>    <int>      <dbl>  
1 N423AS     29  26.89655  
2 N382HA     26  23.80769  
3 N419AS     32  20.56250  
4 N540AA     34  17.26471  
5 N847VA     91  17.09890
```




Combining Rows


bind_rows() bind together two or more datasets by row

Input Datasets



`bind_rows(df1, df2)`

Output Data



Combining Rows cont'd

Task: Add two datasets (with the same variable-names) together by row.


```
df1 <- data_frame(id = 1:2,  
                  name = c("Alice", "Bob"))  
df2 <- data_frame(id = 3:4,  
                  name = c("Charlie", "Dave"))  
  
bind_rows(df1, df2)
```

```
# A tibble: 4 x 2  
  id name  
  <int> <chr>  
1     1 Alice  
2     2   Bob  
3     3 Charlie  
4     4   Dave
```

Combining Columns


`bind_cols()` bind together two or more datasets by column

Input Datasets



`bind_cols(df1, df2)`

Output Data



Combining Columns cont'd

Task: Add two datasets (with the same row-numbers) together by column.

```
df1 <- data_frame(id = 1:2,  
                  name = c("Alice", "Bob"))  
df2 <- data_frame(sales = c(100, 95),  
                  region = c("North", "South"))  
  
bind_cols(df1, df2)
```

```
# A tibble: 2 x 4  
  id name sales region  
  <int> <chr> <dbl> <chr>  
1     1 Alice  100 North  
2     2 Bob    95 South
```

Joins

Combine the datasets by the variable publisher.

“Left” data-frame: superheroes

superhero	alignment	publisher
Batman	good	DC
Joker	bad	DC
Xavier	good	Marvel
Hellboy	good	Dark Horse

“Right” data-frame: address

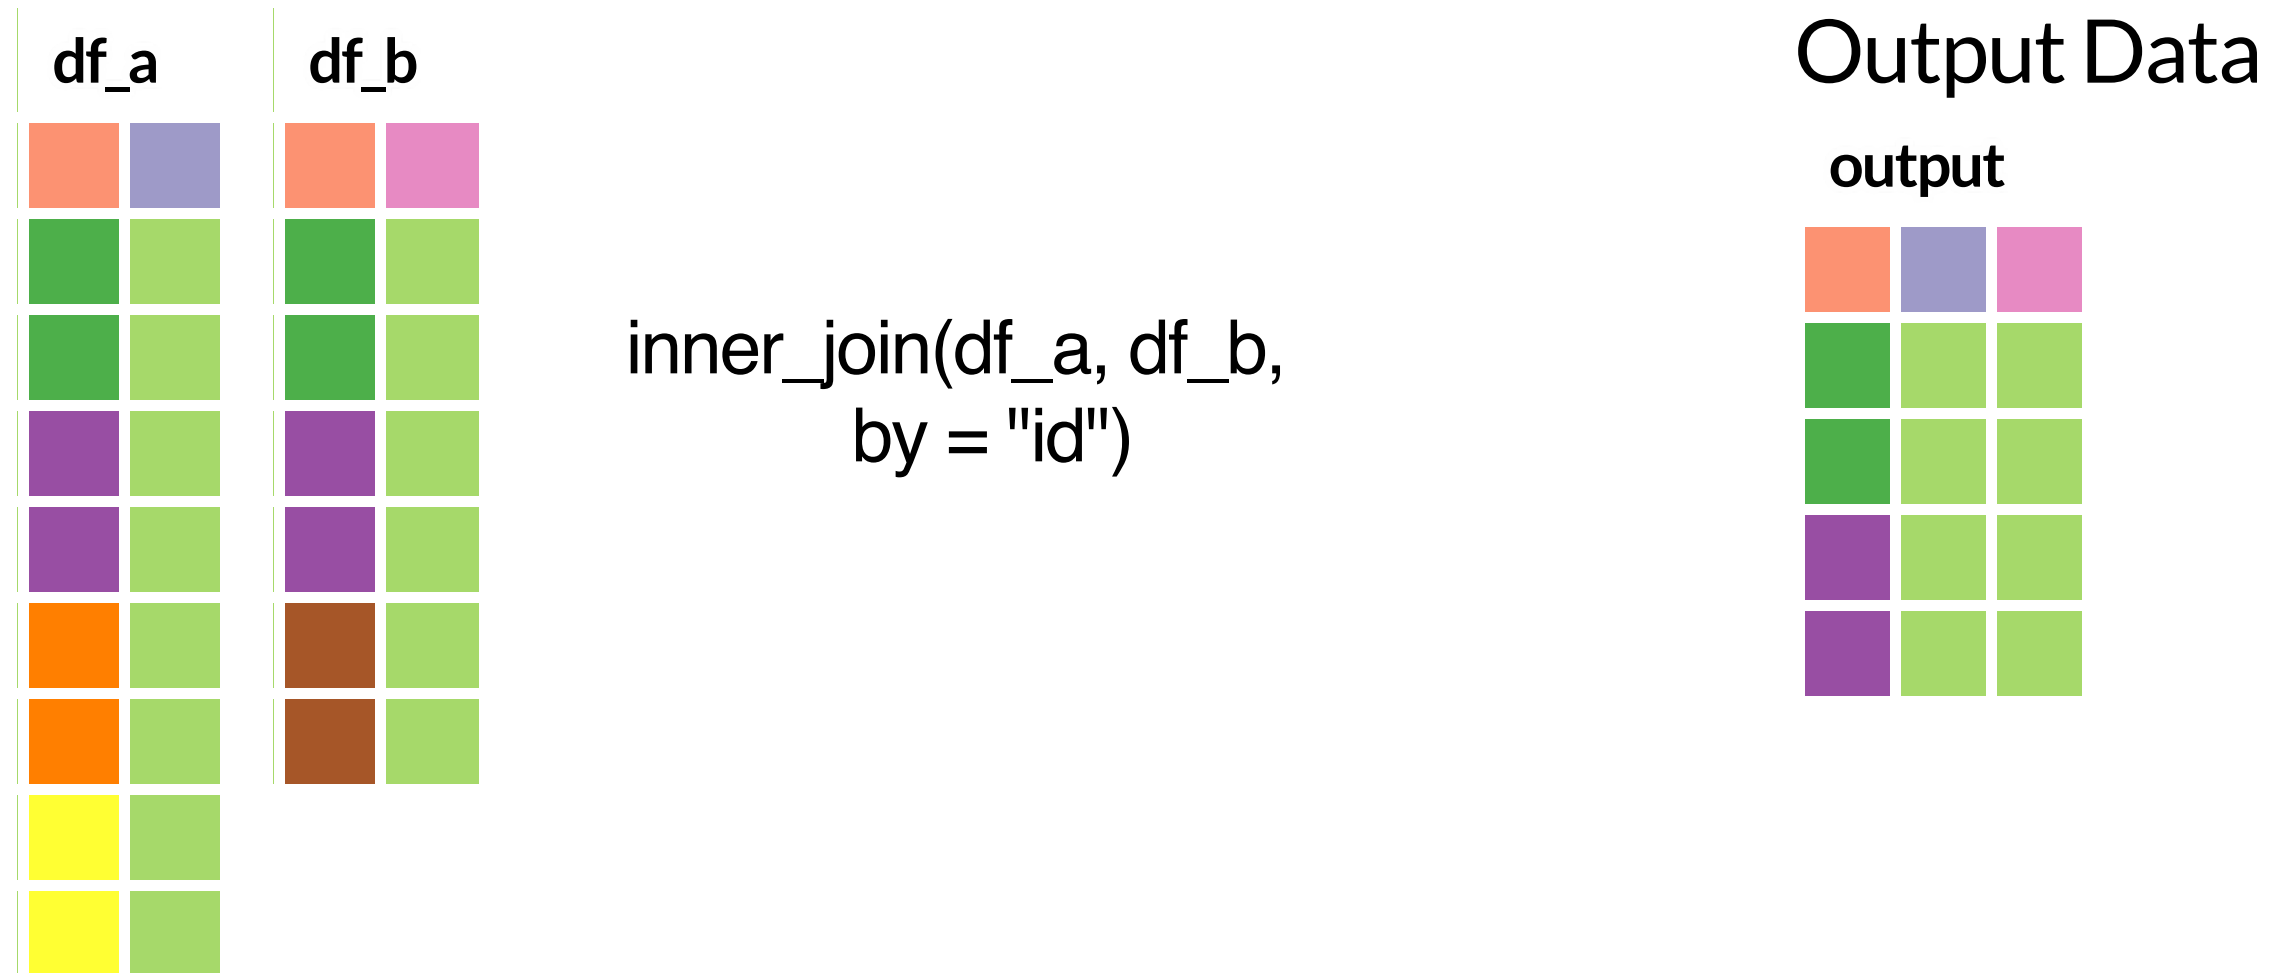
publisher	address
DC	Burbank (CA)
Marvel	NY City (NY)
Image Comics	Berkeley (CA)

Joined data-frame

superhero	alignment	publisher	address
Batman	good	DC	Burbank (CA)
Joker	bad	DC	Burbank (CA)
Xavier	good	Marvel	NY City (NY)
???	???	???	???

Inner Join

`inner_join()` to take only observations found in both datasets



Inner Join cont'd

Task: Find for all superheroes the matching addresses if there is information.

```
# recreate the data
superheroes <- data_frame(
  superhero = c("Batman", "Joker", "Xavier", "Hellboy"),
  alignment = c("good", "bad", "good", "good"),
  publisher = c("DC", "DC", "Marvel", "Dark Horse")
)
address <- data_frame(
  publisher = c("DC", "Marvel", "Image Comics"),
  address   = c("Burbank (CA)", "NY City (NY)",
               "Portland (OR)")
)

# perform the join
inner_join(superheroes, address, by = "publisher")
```

```
# A tibble: 3 x 4
  superhero alignment publisher address
  <chr>    <chr>    <chr>    <chr>
1 Batman   good      DC Burbank (CA)
2 Joker    bad       DC Burbank (CA)
3 Xavier   good     Marvel NY City (NY)
```


Full Join

`full_join()` to take all observations

df_a		df_b	
Orange	Light Purple	Orange	Pink
Green	Light Green	Green	Light Green
Green	Light Green	Green	Light Green
Purple	Light Green	Purple	Light Green
Purple	Light Green	Purple	Light Green
Orange	Light Green	Brown	Light Green
Orange	Light Green	Brown	Light Green
Yellow	Light Green		
Yellow	Light Green		

`full_join(df_a, df_b,
by = "id")`

Output Data

output		
Orange	Light Purple	Pink
Green	Light Green	Light Green
Green	Light Green	Light Green
Purple	Light Green	Light Green
Purple	Light Green	Light Green
Orange	Light Green	Black
Orange	Light Green	Black
Yellow	Light Green	Black
Yellow	Light Green	Black
Brown	Black	Light Green
Brown	Black	Light Green

Full Join cont'd


Task: List all known information about publishing houses and their address.

```
full_join(superheroes, address, by = "publisher")
```

```
# A tibble: 5 x 4
  superhero alignment publisher address
  <chr>    <chr>    <chr>    <chr>
1 Batman   good      DC Burbank (CA)
2 Joker    bad       DC Burbank (CA)
3 Xavier   good      Marvel NY City (NY)
4 Hellboy   good    Dark Horse    <NA>
5 <NA>    <NA> Image Comics Portland (OR)
```

Left Join

left_join() to take all observations of the left dataset



Left Join cont'd

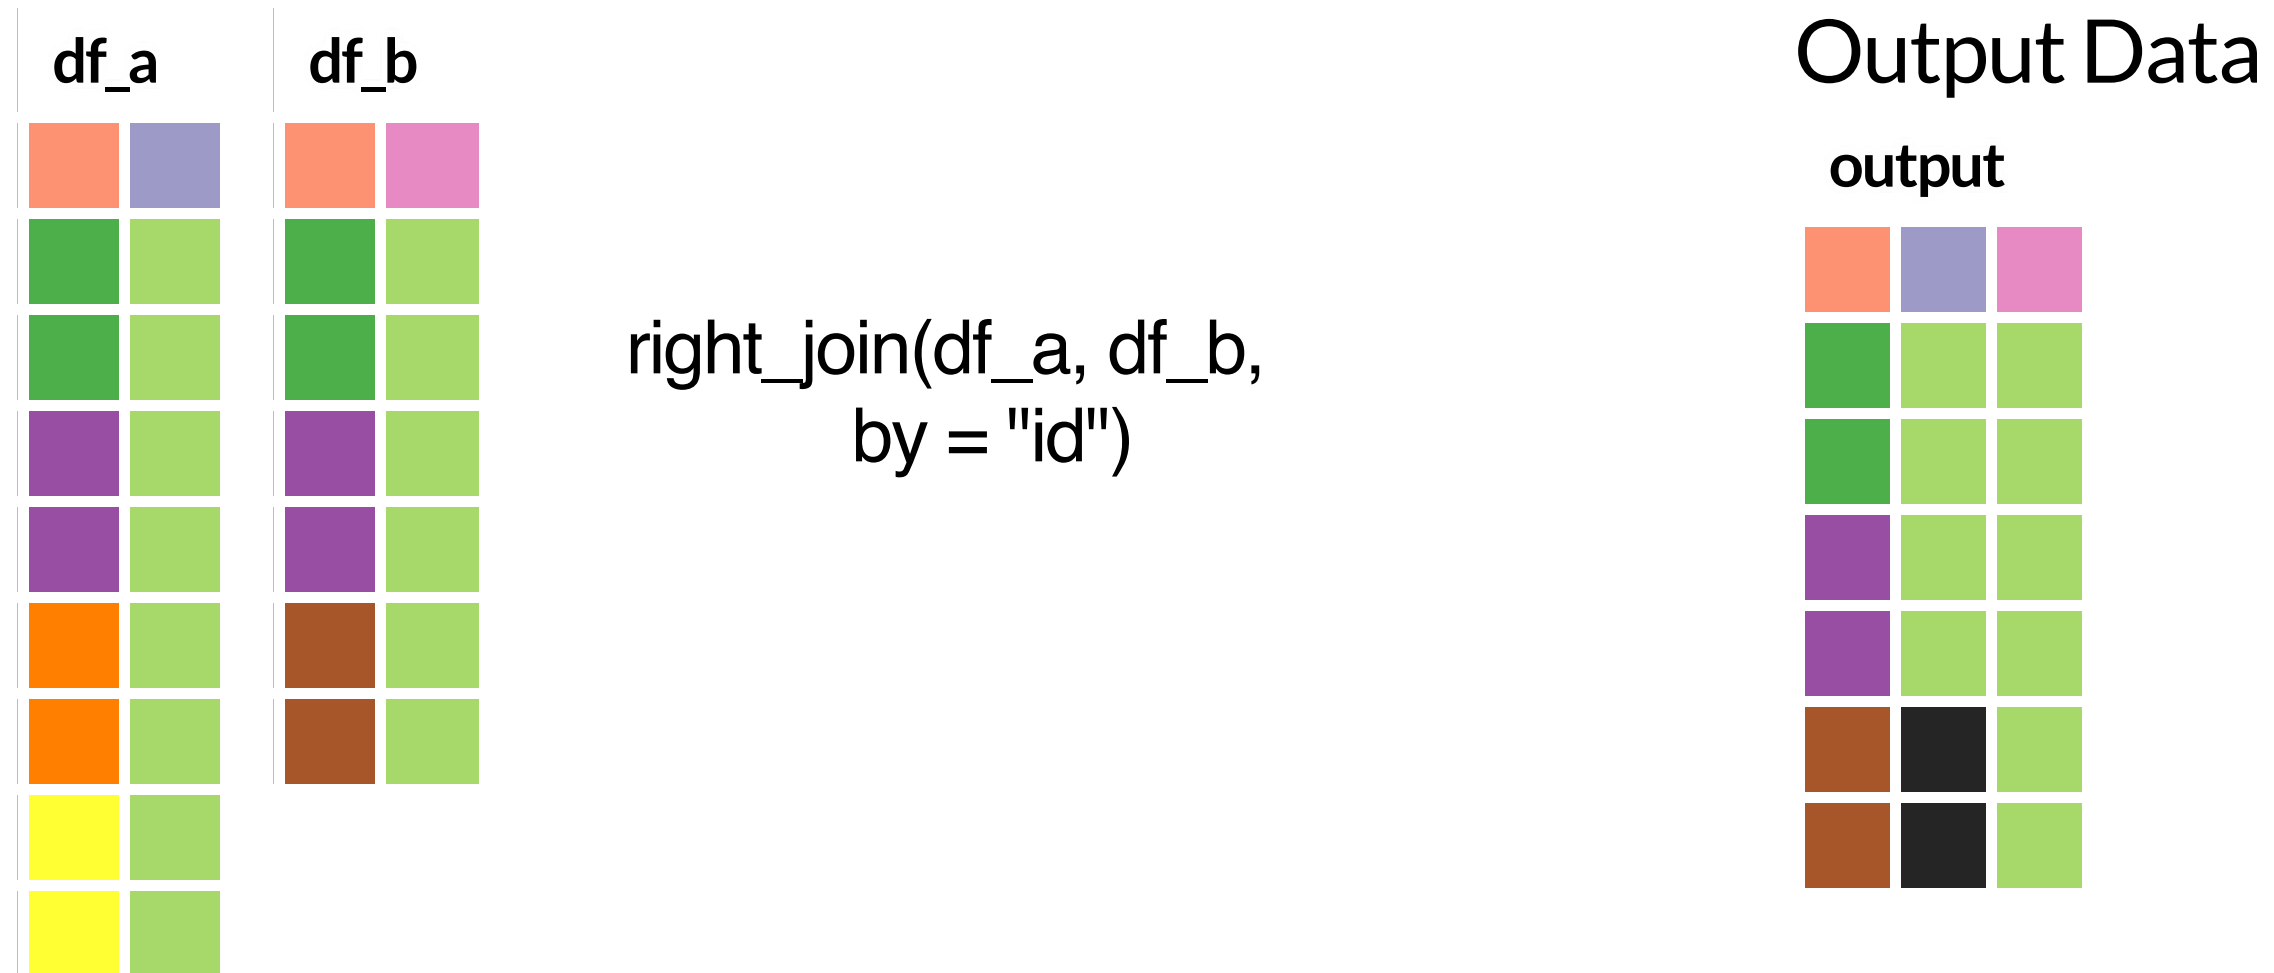
Task: For all superheroes, add their address.

```
left_join(superheroes, address, by = "publisher")
```

```
# A tibble: 4 x 4  
  superhero alignment publisher address  
    <chr>    <chr>    <chr>    <chr>  
1  Batman    good      DC Burbank (CA)  
2   Joker    bad      DC Burbank (CA)  
3  Xavier    good    Marvel NY City (NY)  
4 Hellboy    good Dark Horse    <NA>
```

Right Join

`right_join()` to take all observations of the right dataset



Right Join cont'd


Task: For all publishers, add the superheroes.

```
right_join(superheroes, address, by = "publisher")
```

```
# A tibble: 4 x 4  
  superhero alignment publisher address  
    <chr>    <chr>    <chr>    <chr>  
1  Batman    good      DC Burbank (CA)  
2   Joker    bad      DC Burbank (CA)  
3  Xavier    good    Marvel NY City (NY)  
4   <NA>    <NA> Image Comics Portland (OR)
```

Anti Join

anti_join() to take all observations in df_a that are not in df_b



Anti Join cont'd

Task: Find all superheroes that have no address.

```
anti_join(superheroes, address, by = "publisher")
```


```
# A tibble: 1 x 3  
  superhero alignment publisher  
    <chr>    <chr>    <chr>  
1 Hellboy    good Dark Horse
```

Task: Find all publishers that have no superhero.

```
anti_join(address, superheroes, by = "publisher")
```


```
# A tibble: 1 x 2  
  publisher address  
    <chr>    <chr>  
1 Image Comics Portland (OR)
```


Inner Join
inner_join()




inner_join(df_a, df_b, by = "id")

Outer Join
NA



NA

Full Join
full_join()




full_join(df_a, df_b, by = "id")

Merging Two Datasets

dplyr-syntax


for two data_frames: df_a, df_b

Right Outer Join
anti_join() (reverse)




anti_join(df_b, df_a, by = "id")

Left Join
left_join()




left_join(df_a, df_b, by = "id")

Left Outer Join
anti_join()



anti_join(df_a, df_b, by = "id")

Right Join
right_join()



right_join(df_a, df_b, by = "id")

Additional Resources

- *R for Data Science* by Hadley Wickham Online <http://r4ds.had.co.nz/> or paperback [Amazon](#)
- *Data Wrangling CheatSheet* <https://www.rstudio.com/wp-content/uploads/2015/02/data-wrangling-cheatsheet.pdf>

Questions?

About and Contact

My Ideas: <https://datashenanigan.wordpress.com/>

My Projects: <https://github.com/DavZim/>

My Contact: david_j_zimmermann@hotmail.com

Sources

Images:

- <https://www.tidyverse.org/>
- Rene Magritte: The Treachery of Images 1929
- Ursus Wehrli: The Art of Clean Up 2013
- <http://r4ds.had.co.nz/tidy-data.html>
- <http://fontawesome.io/icon/refresh/>
- <http://www.wallpaperspots.com/fruit-mix-hd-wallpaper/>